

# Vision AI on Nvidia Jetson Nano: Sensor Fusion

## Project Members:

Riya Alias	: 1119262
Nikita Anant Patil	: 1119255
Sreekuttan Sreedharan Thampi	: 1123158
Amrutha Venugopal	: 1119243

# Overview

1. Introduction
2. Flow Diagram
3. Vision AI and Sensor Fusion
4. Nvidia Jetson Nano
5. Edge Impulse SDK
6. Building a dataset in Edge Impulse
7. Designing an Impulse and Validating the model
8. Testing and drawbacks of Edge Impulse
9. YoloV5 model
10. YoloV5 model testing and drawbacks
11. YoloV7 model
12. YoloV7 model testing and results
13. TensorRT model
14. Results
15. Parallel Processing using two cameras
16. Conclusions
17. Future Improvements

# Introduction

- Develop a comprehensive and robust Vision AI system on the NVIDIA Jetson Nano
- Sensor fusion techniques to enhance drowsiness detection and traffic sign detection.
- Drowsiness Detection :
  - Implement a drowsiness detection algorithm that analyzes facial cues
  - Optimize the algorithm for real-time processing using the Jetson Nano's GPU acceleration.
- Traffic Sign Detection:
  - Help drivers by providing real-time information about traffic signals and providing timely alerts or interventions.
  - Implement a mechanism to recognize and react to different types of traffic signs, such as speed limits, stop signs, and yield signs.
- Aims to create an innovative solution that demonstrates the capabilities of sensor fusion in enhancing safety and awareness in driving scenarios.
- The integration of drowsiness detection and traffic sign detection through sensor fusion showcases the power of AI technologies and their potential to positively impact road safety.

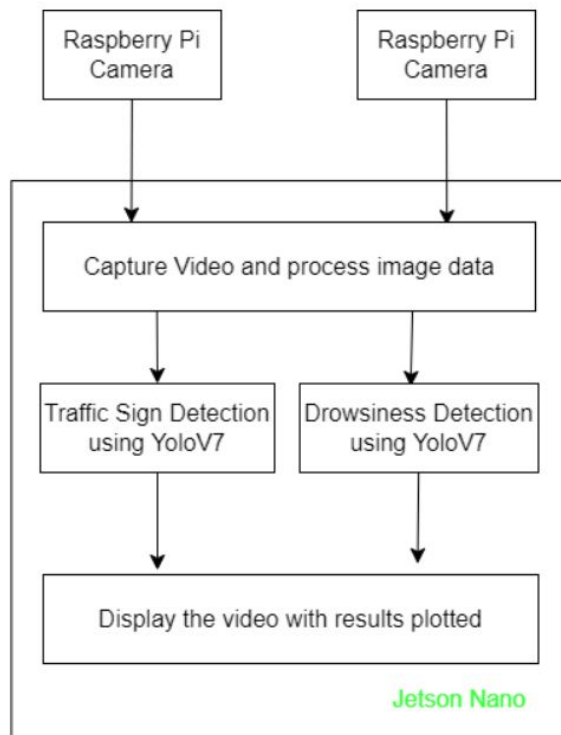
## Driver Monitoring system



## Driver Traffic Signal Assistance



# Flow Diagram



# Vision AI and Sensor Fusion

Combining Vision AI and sensor data is of paramount importance in building intelligent systems that can accurately perceive and understand their environment. This integration offers several key advantages that significantly enhance the capabilities and robustness of various applications.

## **Vision AI :**

- Involves computer vision techniques to enable machines to understand visual information from the world.
- This includes tasks such as image recognition, object detection, semantic segmentation, and even more complex tasks like facial recognition.
- Vision AI leverages deep learning algorithms to analyze visual data and extract meaningful insights

## **Sensor Fusion:**

- Sensor Fusion is the process of integrating data from multiple sensors to create a more accurate and complete representation of the environment.
- Sensors can include cameras, lidar, radar, GPS, IMUs (Inertial Measurement Units), and more.
- By combining data from different sensors, a system gains a better understanding of its surroundings, leading to improved decision-making and situational awareness.

# Nvidia Jetson Nano

- It is a powerful edge computing platform designed to bring the capabilities of artificial intelligence and deep learning to resource-constrained devices at the edge of networks.
- It provides the computational power necessary for running deep learning models and performing complex AI tasks directly on edge devices.
- Nvidia Jetson Nano is exceptionally well-suited for Vision AI and Sensor Fusion applications due to its powerful hardware, dedicated GPU, and AI-focused software support.

## **Key Features:**

- GPU Acceleration
- Quad-Core CPU
- Memory and Storage
- Connectivity
- AI Software Support

# Hardware

- Uses Nvidia Jetson Nano Kit to develop a system that can perform real time image processing and object detection using multiple cameras.
- **Jetson Nano:**
  - small, powerful computer
  - Allows to run multiple neural networks in parallel
  - Applications: image classification, object detection, segmentation, and speech processing.
  - GPU : NVIDIA Maxwell™ architecture with 128 NVIDIA CUDA® cores
  - CPU : Quad-core ARM® Cortex®-A57 MPCore processor
  - Memory : 4 GB 64-bit LPDDR4
- **Raspberry Pi Camera**
  - Has a Sony IMX219 8-megapixel sensor.
  - Can be used to take high-definition video, as well as stills photographs.
  - supports 1080p30, 720p60 and VGA90 video modes

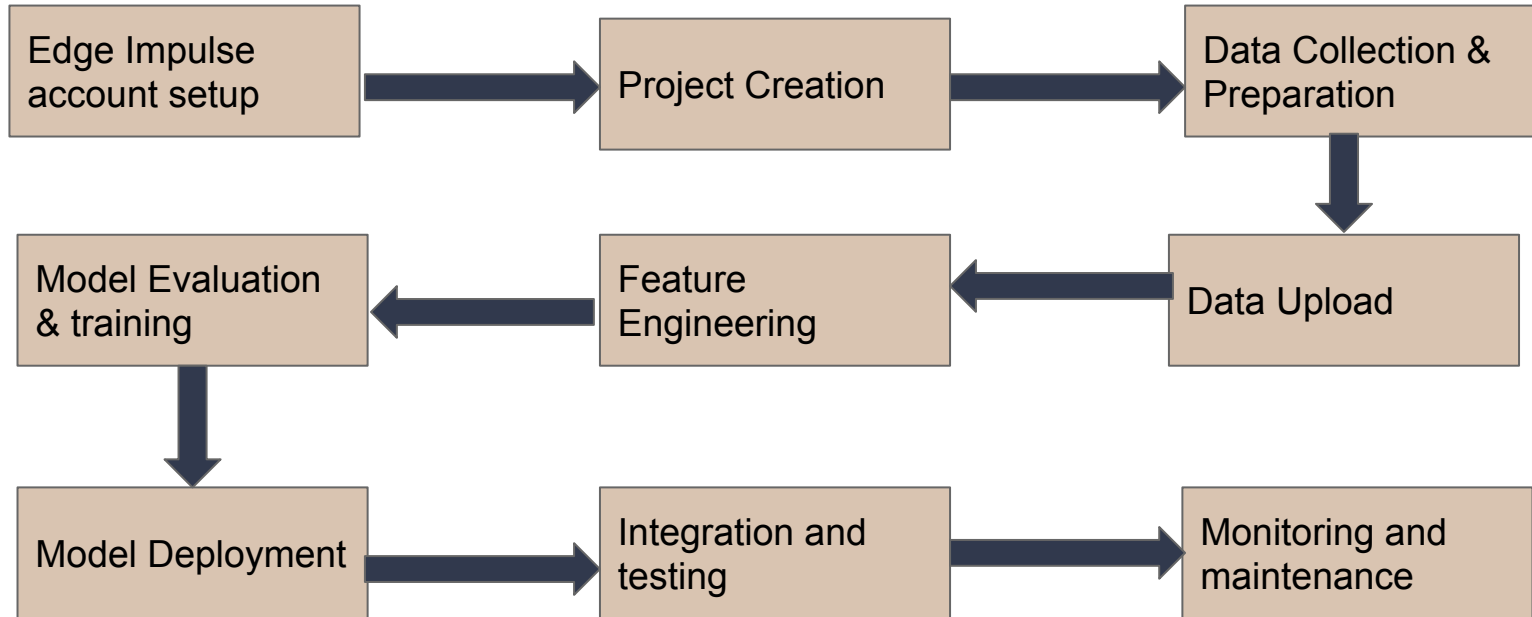


# Implementation

# Edge Impulse

- A suite of engineering tools for building, training, and deploying machine learning models on edge devices.
- Cloud based platform for machine learning on edge devices.
- No Coding Required: provides a user-friendly interface
- Edge Computing: ensuring low latency and real-time processing.
- Support for Various Sensors
- Jetson Nano is fully supported by Edge Impulse
- Supports various ML tools like PyTorch and TensorFlow

# Flow Diagram



# Building a Dataset : data acquisition

The screenshot displays the Edge Impulse web interface. On the left is a sidebar with navigation links: Dashboard, Devices, Data acquisition, Impulse design, and FON Trainer. Below these is a 'Try Enterprise Free' banner. The main content area has a top navigation bar with 'Dataset', 'Data sources', and 'Labeling queue (5)'. Below this, there are two summary cards: 'DATA COLLEC... 335 items' with a red donut chart and 'TRAIN / TES... 81% ...' with a green donut chart. A 'Collect data' section on the right includes a link to 'Connect a device'. Below the summary cards is a 'Dataset' table with columns for 'SAMPLE NAME', 'LABELS', and 'ADDED'. The table lists five samples, all labeled 'Non\_drowsy'. A dark blue overlay on the right side of the table says 'Click on a sample to load...'. The user's name 'Riya / Sample' is in the top right corner.

**EDGE IMPULSE**

Dashboard  
Devices  
Data acquisition  
Impulse design  
FON Trainer

**Try Enterprise Free**  
Get access to high job limits and training on GPUs.  
[Start free trial](#)

Riya / Sample

**Dataset** | Data sources | Labeling queue (5)

DATA COLLEC...  
335 items

TRAIN / TES...  
81% ...

**Collect data**  
[Connect a device](#) to start building your dataset.

**Dataset**

Training (266) | Test (69)

SAMPLE NAME	LABELS	ADDED
2623	Non_drowsy	Jun 01 2023, 1...
2621	Non_drowsy	Jun 01 2023, 1...
2582	Non_drowsy	Jun 01 2023, 1...
2606	Non_drowsy	Jun 01 2023, 1...
2590	Non_drowsy	Jun 01 2023, 1...

**RAW DATA**  
Click on a sample to load...

# Designing an Impulse

An impulse takes the raw data, adjusts the image size, uses a preprocessing block to manipulate the image, and then uses a learning block to classify new data.

The interface displays a workflow for designing an impulse, consisting of four main blocks arranged horizontally:

- Image data** (Red block):
  - Input axes: image
  - Image width: 96
  - Image height: 96
  - Resize mode: Fit shortest axis (dropdown menu)
  - Icon: A person icon with a red arrow pointing to a square icon.
  - Footer: For object detection use a square image size, e.g. 96x96, 160x160 or 320x320.
- Image** (White block):
  - Name: Image
  - Input axes (1):
    - ☒ image
- Object Detection (Images)** (Blue block):
  - Name: Object detection
  - Input features:
    - ☒ Image
  - Output features: 2 (Non\_drowsy, drowsy)
- Output features** (Green block):
  - 2 (Non\_drowsy, drowsy)

Below the blocks are two dashed boxes for adding new blocks:

- Add a processing block** (Light blue dashed box with a lightning bolt icon)
- Add a learning block** (Light blue dashed box with a flask icon)

A **Save Impulse** button is located to the right of the Output features block.

# Configuring the transfer learning block

- Start Training the neural network model
- Take the image data as an input, and try to map this to one of the two classes.
- Transfer Learning : only retraining the upper layers of a neural network, leading to much more reliable models that train in a fraction of the time and work with substantially smaller datasets.

The screenshot displays a user interface for configuring and training a neural network model. The interface is divided into two main panels: 'Neural Network settings' on the left and 'Training output' on the right.

**Neural Network settings:**

- Training settings:**
  - Number of training cycles: 100
  - Learning rate: 0.005
  - Data augmentation: ☒
- Advanced training settings:** (Collapsed)
- Neural network architecture:**
  - Input layer (27,648 features)
  - Model: FOMO (Faster Objects, More Objects) MobileNetV2 0.35
  - Choose a different model

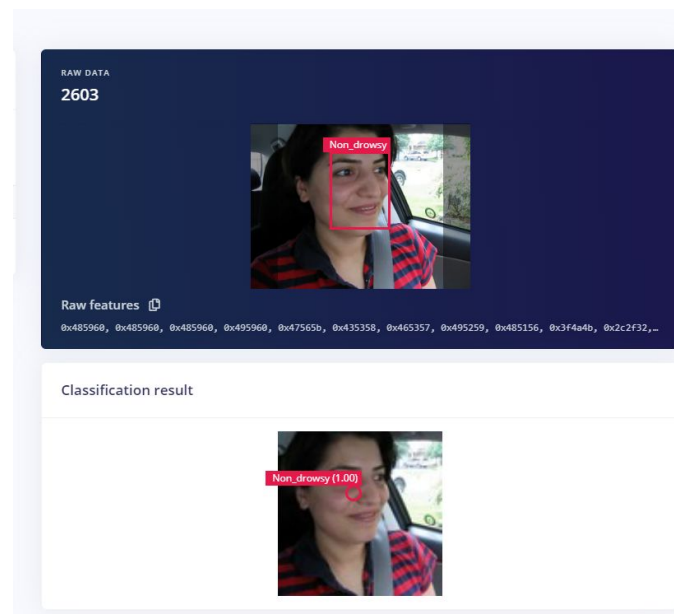
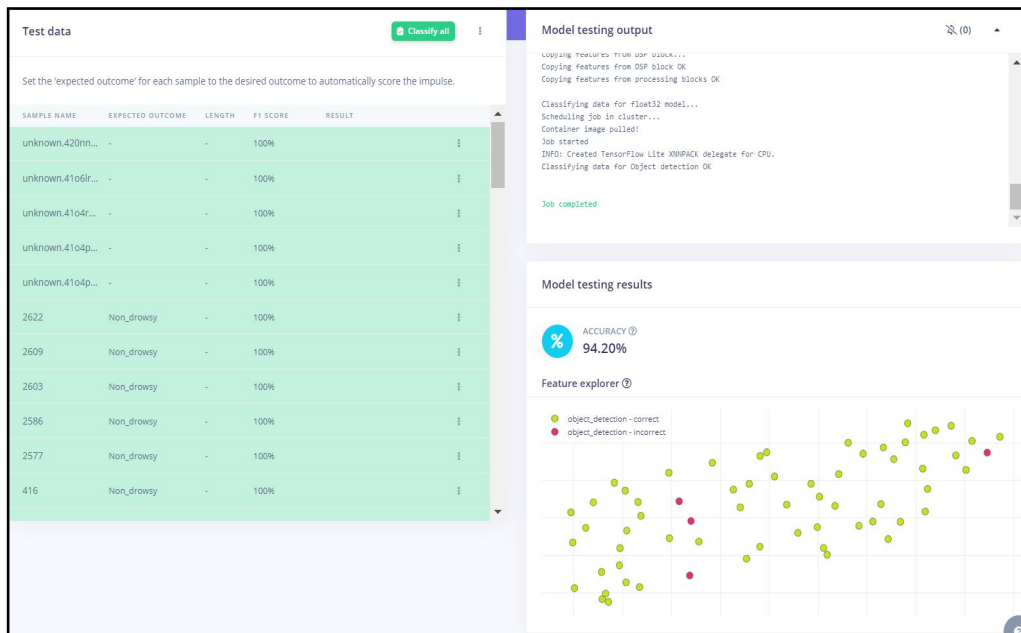
**Training output:**

- Model:** Quantized (int8)
- Last training performance (validation set):**
  - F1 SCORE: 85.5%
- Confusion matrix (validation set):**

	BACKGROUND	NON_DROWSY	DROWSY
BACKGROUND	99.9%	0.1%	0.1%
NON_DROWSY	9.4%	90.6%	0%
DROWSY	18.2%	0%	81.8%
F1 SCORE	1.00	0.88	0.82
- On-device performance:**
  - INFERRING TIME: 4 ms.
  - PEAK RAM USAGE: 239.1K
  - FLASH USAGE: 72.8K

# Validating the model

Use the data in the testing dataset to validate how well the model will work in the real world.



# Testing and Drawbacks of Edge impulse

- Model Performance:  
We were able to achieve an accuracy of 94.20 with the test dataset.  
But with new real time images, the detection results were not that good.
- Model Training Constraints:  
While Edge Impulse provides tools for model training, the training process might be slower and less versatile compared to using more robust cloud-based training platforms.
- Edge Impulse provides two different methods to perform object detection:  
Using MobileNetV2 SSD FPN  
Using FOMO



# YoloV5

- YoloV5 is an object detection model, designed to deliver high-speed, high-accuracy results in real-time.
- It is built on PyTorch

## **Implementation Steps:**

### **Step 1: Dataset Collection and preparation**

- For Drowsiness Detection : Dataset collected via camera (custom dataset)  
Classes 2 : drowsy and awake  
Number of images : 566
- For Traffic Sign Detection : Dataset took from German Traffic Sign Detection Benchmark (GTSDB)  
Classes 4 : prohibitory, mandatory, danger and other  
Number of images : 900

Divided into two group for training and validating 80-20.

Label text file contains the class, height and width of the image

## Step 2 : Setting up the environment in Jupyter Notebook

- Cloning the YoloV5 repository and installing all the necessary packages like opencv, numpy, torch, torchvision etc required to train our model.

## Step 3: Train and validate the model

```
!cd yolov5 && python train.py --img 320 --batch 16 --epochs 50 --data dataset.yml --weights yolov5s.pt --workers 2
```

8, 1.44s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	77%	#####7	17/22	[00:24<00:0
7, 1.42s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	82%	#####1	18/22	[00:25<00:0
5, 1.42s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	86%	#####6	19/22	[00:27<00:0
4, 1.41s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	91%	#####	20/22	[00:28<00:0
2, 1.41s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	95%	#####5	21/22	[00:29<00:0
1, 1.41s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	100%	#####	22/22	[00:30<00:0
0, 1.08s/it]	Class	Images	Instances	P	R	mAP50	mAP50-95:	100%	#####	22/22	[00:30<00:0
0, 1.37s/it]											
	all	678	113	0.165	0.99	0.205	0.16				
	drowsy	678	50	0.149	0.98	0.201	0.174				
	awake	678	63	0.18	1	0.209	0.147				

Results saved to runs\train\exp8

Training is done using train.py

Image size : 320\* 320

Batch : 16

Epochs : 50

# YoloV5 model test results and drawbacks

Drawback of using Jupyter Notebook for training:

- The training took approximately 36 hours for traffic sign detection(900 images). Hence could not be used for large number of datasets

Drawback of using YoloV5

- Less prediction rates



# YoloV7

- YOLOv7 provides a faster and stronger network architecture
- Provides a more effective feature integration method, more accurate object detection performance, a more robust loss function.

## Comparison with YoloV5

- Higher inference speed (more than 20 FPS Faster).
- YoloV7 reduces the number of parameters by 22%, require 8% less computation power and has an increased average precision of 2.2%

# Training YoloV7 using Google Collab

## Implementation Steps:

### Step 1: Dataset Collection and preparation (used the same dataset)

- For Drowsiness Detection : Dataset collected via camera (custom dataset)  
Classes 2 : drowsy and awake  
Number of images : 800
- For Traffic Sign Detection : Dataset took from German Traffic Sign Detection Benchmark (GTSDDB)  
Classes 4 : prohibitory, mandatory, danger and other  
Number of images : 900

### Step 2: Clone the repository for YoloV7

### Step 3: Train and validate the model

```
!python train.py --device 0 --batch-size 16 --epochs 30 --img 640 640 --data data/custom_data.yaml --hyp data/hyp.scratch.custom.yaml --cfg cfg/training/yolov7x-custom.yaml --weights yolov7x.pt --name yolov7x-custom
```

2023-08-08 11:05:53.280887: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-08-08 11:05:54.599694: W tensorflow/compiler/tf2tensorrt/utils/py\_utils.cc:38] TF-TRT Warning: Could not find TensorRT

YOLOR 🚀 v0.1-126-g84932d7 torch 2.0.1+cu118 CUDA:0 (Tesla T4, 15101.8125MB)

Namespace(weights='yolov7x.pt', cfg='cfg/training/yolov7x-custom.yaml', data='data/custom\_data.yaml', hyp='data/hyp.scratch.custom.yaml', epochs=30, batch\_size=16, img\_size=[640, 640], rect=False, resume=False, nosave=False)

**tensorboard:** Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>

**hyperparameters:** lr=0.01, lrf=0.1, momentum=0.937, weight\_decay=0.0005, warmup\_epochs=3.0, warmup\_momentum=0.8, warmup\_bias\_lr=0.1, box=0.05, cls=0.3, cls\_pw=1.0, obj=0.7, obj\_pw=1.0, iou\_t=0.2, anchor\_t=4.0, fl\_gamma=0.0

**wandb:** Install Weights & Biases for YOLOR logging with 'pip install wandb' (recommended)

Image size : 640\*640

Batch : 16

Epochs :30

Model : yolov7-x model

Used GPU for training

Testing in Google Colab: Tested images and observed that the prediction is now 80%-90%.

# Testing and Analysis of YoloV7 model

## **Step 4: Deployment into Jetson Nano**

- Created separate virtual environments for both the Yolo models for testing
- Installed required packages like OpenCV, Numpy, Matplotlib, etc
- Installed Jupyter Notebook
- Cloned the repository for YoloV7

## **Step 5: Testing the model using images and videos in Jetson Nano**

- Tested few images and videos and observed that the inference time is not as expected.
- The processing speed and the FPS was low.



YoloV7 Prediction results



# TensorRT

- High performance deep learning inference runtime for image classification, segmentation, and object detection neural networks.
- Built on NVIDIA's parallel programming model : CUDA

## Steps for Converting the YoloV7 model to TensorRT model

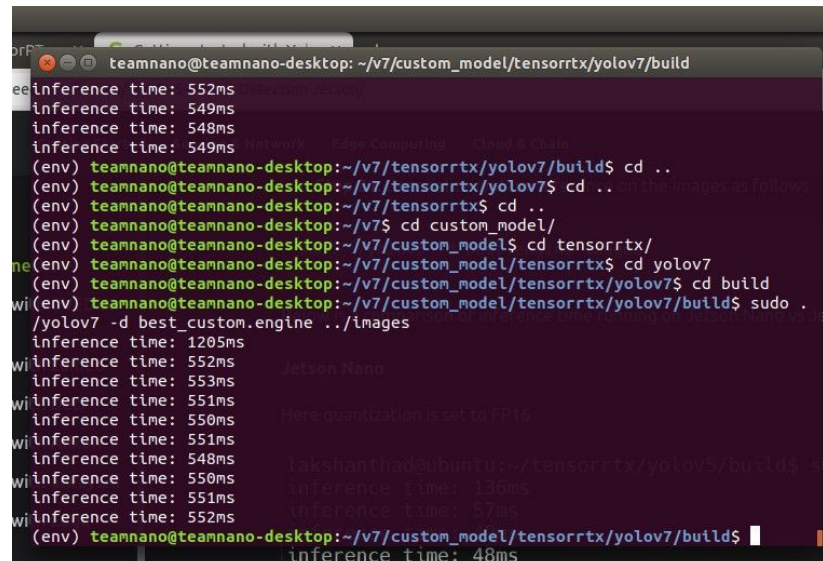
- Cloned the TensorRT repository to the system.

```
git clone https://github.com/wang-xinyu/tensorrtx
```

- Converted the .pt weights of the YoloV7 to the .wts file using the gen\_wts.py.
- The generated .wts file is then converted to .engine file. This engine file is used for the testing and detections.


# Testing and Analysis

- Tested the images and videos using the generated engine file.
- Observed that the inference time is now increased to 2FPS.



```
teamnano@teamnano-desktop: ~/v7/custom_model/tensorrtx/yolov7/build
inference time: 552ms
inference time: 549ms
inference time: 548ms
inference time: 549ms
(env) teamnano@teamnano-desktop:~/v7/tensorrtx/yolov7/build$ cd ..
(env) teamnano@teamnano-desktop:~/v7/tensorrtx/yolov7$ cd .. in the images as follows
(env) teamnano@teamnano-desktop:~/v7/tensorrtx$ cd ..
(env) teamnano@teamnano-desktop:~/v7$ cd custom_model/
(env) teamnano@teamnano-desktop:~/v7/custom_model$ cd tensorrtx/
(env) teamnano@teamnano-desktop:~/v7/custom_model/tensorrtx$ cd yolov7
(env) teamnano@teamnano-desktop:~/v7/custom_model/tensorrtx/yolov7$ cd build
wi (env) teamnano@teamnano-desktop:~/v7/custom_model/tensorrtx/yolov7/build$ sudo .
/v7/yolov7 -d best_custom.engine ../images
inference time: 1205ms
wi inference time: 552ms
inference time: 553ms
wi inference time: 551ms
inference time: 550ms
wi inference time: 551ms
inference time: 548ms
wi inference time: 550ms
inference time: 551ms
inference time: 552ms
(env) teamnano@teamnano-desktop:~/v7/custom_model/tensorrtx/yolov7/build$ inference time: 48ms
```

# Results



Video testing




Image testing

# Results

# Video Testing



# Real Time Testing





# Traffic Sign detection – sample predicted output



# Parallel Processing of two cameras

- For testing both the models (driver assistance system and driver monitoring system), live feed from both the cameras are taken and processed.
- Multithreading is used to run the inference on each camera feed independently.
- We observed that running resource-intensive tasks concurrently on the Jetson Nano led to power consumption levels that exceeded the device's capacity.
- As a result, the system experienced performance throttling to prevent overheating and maintain power efficiency.
- Given the limited power available, the execution of parallel tasks placed excessive demand on the device's resources, caused unexpected slowdowns, degraded inference performance, and an automatic rebooting.



# Conclusion

- By integrating cutting-edge AI technologies, YOLOv7 and TensorRT, we've created a robust solution that addresses critical aspects of road safety.
- We faced many challenges during the implementation and were able to resolve most of them, including the jetson nano rebooting issues,
- Learned on how to train the model on custom datasets and testing them.
- Learned about different deep learning algorithms.
- Learned how to deploy and test the model in Jetson Nano.
- In conclusion, our project demonstrates the impact of AI-driven driver assistance systems. By leveraging YOLOv7 and TensorRT, we've crafted a technology-driven solution that has the potential to significantly reduce accidents and create safer driving environments.

# Future Improvements

## **Drowsiness Detection**

1. Real-time Feedback and Alerts:
  - Alerting the driver, such as haptic feedback on the steering wheel or seat.
2. Integration with Autonomous Systems:
  - Integrate drowsiness detection into autonomous vehicles to ensure that the system can take control when the driver becomes drowsy.

## **Traffic Sign Detection**

1. Multi-Sign Detection:
  - Extend the system to detect multiple signs in a single frame, allowing for a more comprehensive understanding of the traffic environment.
2. Real-time Traffic Analysis:
  - Integrate the detected traffic sign data with mapping and navigation systems to provide real-time information to the driver about upcoming road conditions and regulations.

# Future Improvements

## **Parallel Processing** (Sensor fusion) improvement using **DeepStream SDKs**:

- For better performance and efficiency, consider using NVIDIA's DeepStream SDK, which is designed for real-time AI inference on Jetson devices and includes features for working with multiple cameras and models.
- Additionally, optimizing the models and algorithms for efficiency and resource utilization may help mitigate the effects of power throttling which hindered the parallel processing ability.

# Reference

- <https://github.com/nicknochnack/YOLO-Drowsiness-Detection>
- <https://github.com/AarohiSingla/yolov5>
- <https://www.kaggle.com/datasets/valentynsichkar/traffic-signs-dataset-in-yolo-format>
- <https://wiki.seedstudio.com/YOLOv5-Object-Detection-Jetson/>
- <https://medium.com/@jurespeh/yolov7-with-tensorrt-on-jetson-nano-with-python-script-example-63099fa7c8a5>
- <https://developer.nvidia.com/embedded/learn/get-started-jetson-nano-devkit>
- <https://towardsdatascience.com/applied-deep-learning-part-1-artificial-neural-networks-d7834f67a4f6>
- <https://viso.ai/deep-learning/yolov7-guide/>
- <https://github.com/WongKinYiu/yolov7>
- <https://docs.edgeimpulse.com/docs/tutorials/end-to-end-tutorials/object-detection/object-detection>
- <https://github.com/ultralytics/yolov5>
- <https://github.com/JetsonHacksNano/CSI-Camera>
- <https://www.kaggle.com/datasets/valentynsichkar/traffic-signs-dataset-in-yolo-format>

Thank You!